



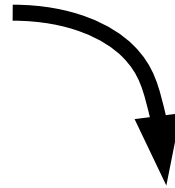



# Advanced Scikit-Learn

Andreas Mueller (NYU Center for Data Science, scikit-learn)



# Me





**agramfort**  
Alexandre Gramfort



**AlexanderFabisch**  
Alexander Fabisch



**alextp**  
Alexandre Passos



**amueller**  
Andreas Mueller



**arjoly**  
Arnaud Joly



**bdholt1**  
Brian Holt



**GaelVaroquaux**  
Gael Varoquaux



**glouppe**  
Gilles Louppe



**jakevdp**  
Jake Vanderplas



**jaquesgrobler**  
Jaques Grobler



**jnothman**



**kastnerkyle**  
Kyle Kastner



**bthirion**  
bthirion



**chrisfilo**  
Chris Filo Gorgole...



**cournape**  
David Courapeau



**duchesnay**  
Duchesnay



**dwf**  
David Warde-Farley



**fabianp**  
Fabian Pedregosa



**kuantkid**  
Wei Li



**larsmans**  
Lars



**lucidfrontier45**  
Shiqiao Du



**mblondel**  
Mathieu Blondel



**MechCoder**  
Manoj Kumar



**ndawe**  
Noel Dawe



**NelleV**  
Varoquaux



**ogrisel**  
Olivier Grisel



**paolo-losi**  
Paolo Losi



**pprett**  
Peter Prettenhofer



**robertlayton**  
Robert Layton



**ronw**  
Ron Weiss



**satra**  
Satrajit Ghosh



**sklearn-ci**



**vene**  
Vlad Niculae



**VirgileFritsch**  
Virgile Fritsch



**vmichel**  
Vincent Michel



**yarikoptic**  
Yaroslav Halchenko



# Representing Data

$$X = \begin{pmatrix} 1.1 & 2.2 & 3.4 & 5.6 & 1.0 \\ 6.7 & 0.5 & 0.4 & 2.6 & 1.6 \\ 2.4 & 9.3 & 7.3 & 6.4 & 2.8 \\ 1.5 & 0.0 & 4.3 & 8.3 & 3.4 \\ 0.5 & 3.5 & 8.1 & 3.6 & 4.6 \\ 5.1 & 9.7 & 3.5 & 7.9 & 5.1 \\ 3.7 & 7.8 & 2.6 & 3.2 & 6.3 \end{pmatrix}$$

# Representing Data

one sample

$X =$

1.1	2.2	3.4	5.6	1.0
6.7	0.5	0.4	2.6	1.6
2.4	9.3	7.3	6.4	2.8
1.5	0.0	4.3	8.3	3.4
0.5	3.5	8.1	3.6	4.6
5.1	9.7	3.5	7.9	5.1
3.7	7.8	2.6	3.2	6.3

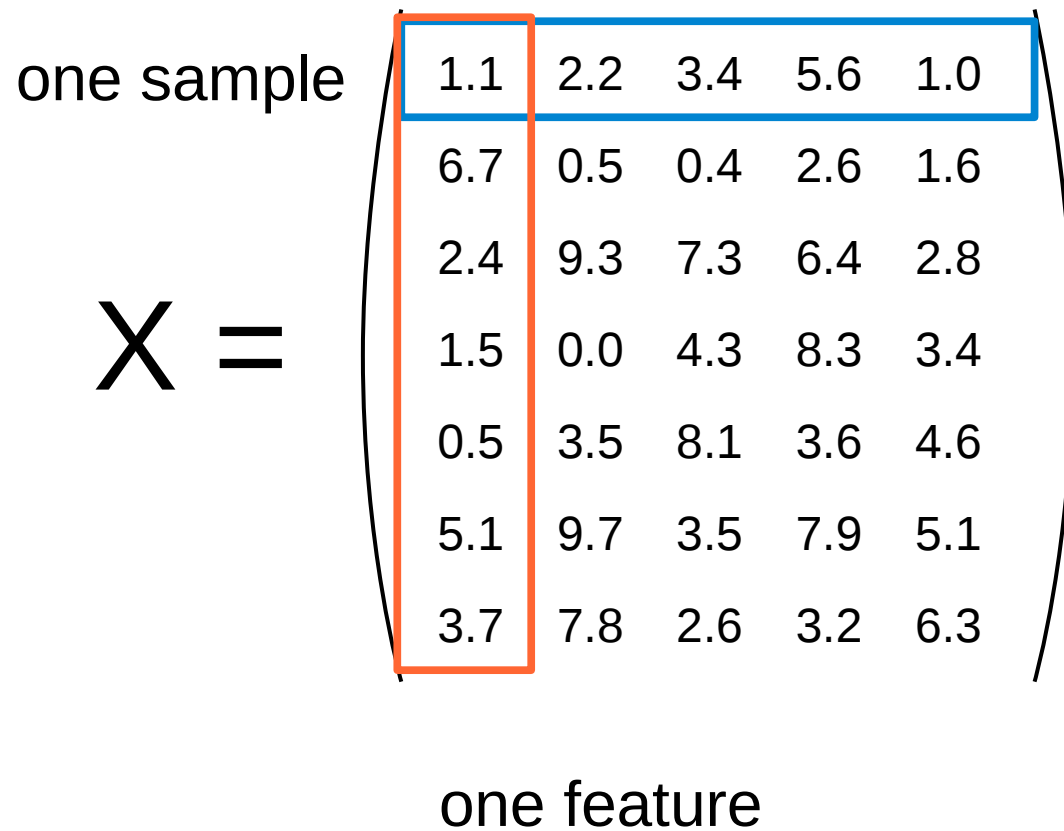
# Representing Data

one sample

$X =$

1.1	2.2	3.4	5.6	1.0
6.7	0.5	0.4	2.6	1.6
2.4	9.3	7.3	6.4	2.8
1.5	0.0	4.3	8.3	3.4
0.5	3.5	8.1	3.6	4.6
5.1	9.7	3.5	7.9	5.1
3.7	7.8	2.6	3.2	6.3

one feature



# Representing Data

one sample

$$X = \begin{pmatrix} 1.1 & 2.2 & 3.4 & 5.6 & 1.0 \\ 6.7 & 0.5 & 0.4 & 2.6 & 1.6 \\ 2.4 & 9.3 & 7.3 & 6.4 & 2.8 \\ 1.5 & 0.0 & 4.3 & 8.3 & 3.4 \\ 0.5 & 3.5 & 8.1 & 3.6 & 4.6 \\ 5.1 & 9.7 & 3.5 & 7.9 & 5.1 \\ 3.7 & 7.8 & 2.6 & 3.2 & 6.3 \end{pmatrix}$$

one feature

$$y = \begin{pmatrix} 1.6 \\ 2.7 \\ 4.4 \\ 0.5 \\ 0.2 \\ 5.6 \\ 6.7 \end{pmatrix}$$

outputs / labels

# Basic API

```
estimator.fit(X, [y])
```

```
estimator.predict
```

```
estimator.transform
```

---

Classification

Preprocessing

Regression

Dimensionality reduction

Clustering


Feature selection

Feature extraction




# Model selection and model complexity (aka bias-variance tradeoff)


# Overfitting and Underfitting



# Overfitting and Underfitting



# Overfitting and Underfitting



# Cross-Validation

```
from sklearn.cross_validation import cross_val_score  
  
scores = cross_val_score(SVC(), X, y, cv=5)  
print(scores)  
  
>> [ 0.92  1.    1.    1.    1. ]
```

# Cross-Validation

```
from sklearn.cross_validation import cross_val_score

scores = cross_val_score(SVC(), X, y, cv=5)
print(scores)

>> [ 0.92  1.    1.    1.    1. ]

cv_ss = ShuffleSplit(len(X_train), test_size=.3,
                    n_iter=10)
scores_shuffle_split = cross_val_score(SVC(), X, y,
                                       cv=cv_ss)
```

# Cross-Validation

```
from sklearn.cross_validation import cross_val_score

scores = cross_val_score(SVC(), X, y, cv=5)
print(scores)

>> [ 0.92  1.    1.    1.    1. ]

cv_ss = ShuffleSplit(len(X_train), test_size=.3,
                    n_iter=10)
scores_shuffle_split = cross_val_score(SVC(), X, y,
                                       cv=cv_ss)

cv_labels = LabelKFold(labels, n_folds=5)
scores_pout = cross_val_score(SVC(), X, y, cv=cv_labels)
```

# Cross-Validation

```
from sklearn.cross_validation import cross_val_score

scores = cross_val_score(SVC(), X, y, cv=5)
print(scores)

>> [ 0.92  1.    1.    1.    1. ]

cv_ss = ShuffleSplit(len(X_train), test_size=.3,
                    n_iter=10)
scores_shuffle_split = cross_val_score(SVC(), X, y,
                                       cv=cv_ss)

cv_labels = LabelKFold(labels, n_folds=5)
scores_pout = cross_val_score(SVC(), X, y, cv=cv_labels)
```



All Data

Training data

Test data

All Data

Training data      Test data

Fold 1   Fold 2   Fold 3   Fold 4   Fold 5

Split 1   Fold 1   Fold 2   Fold 3   Fold 4   Fold 5

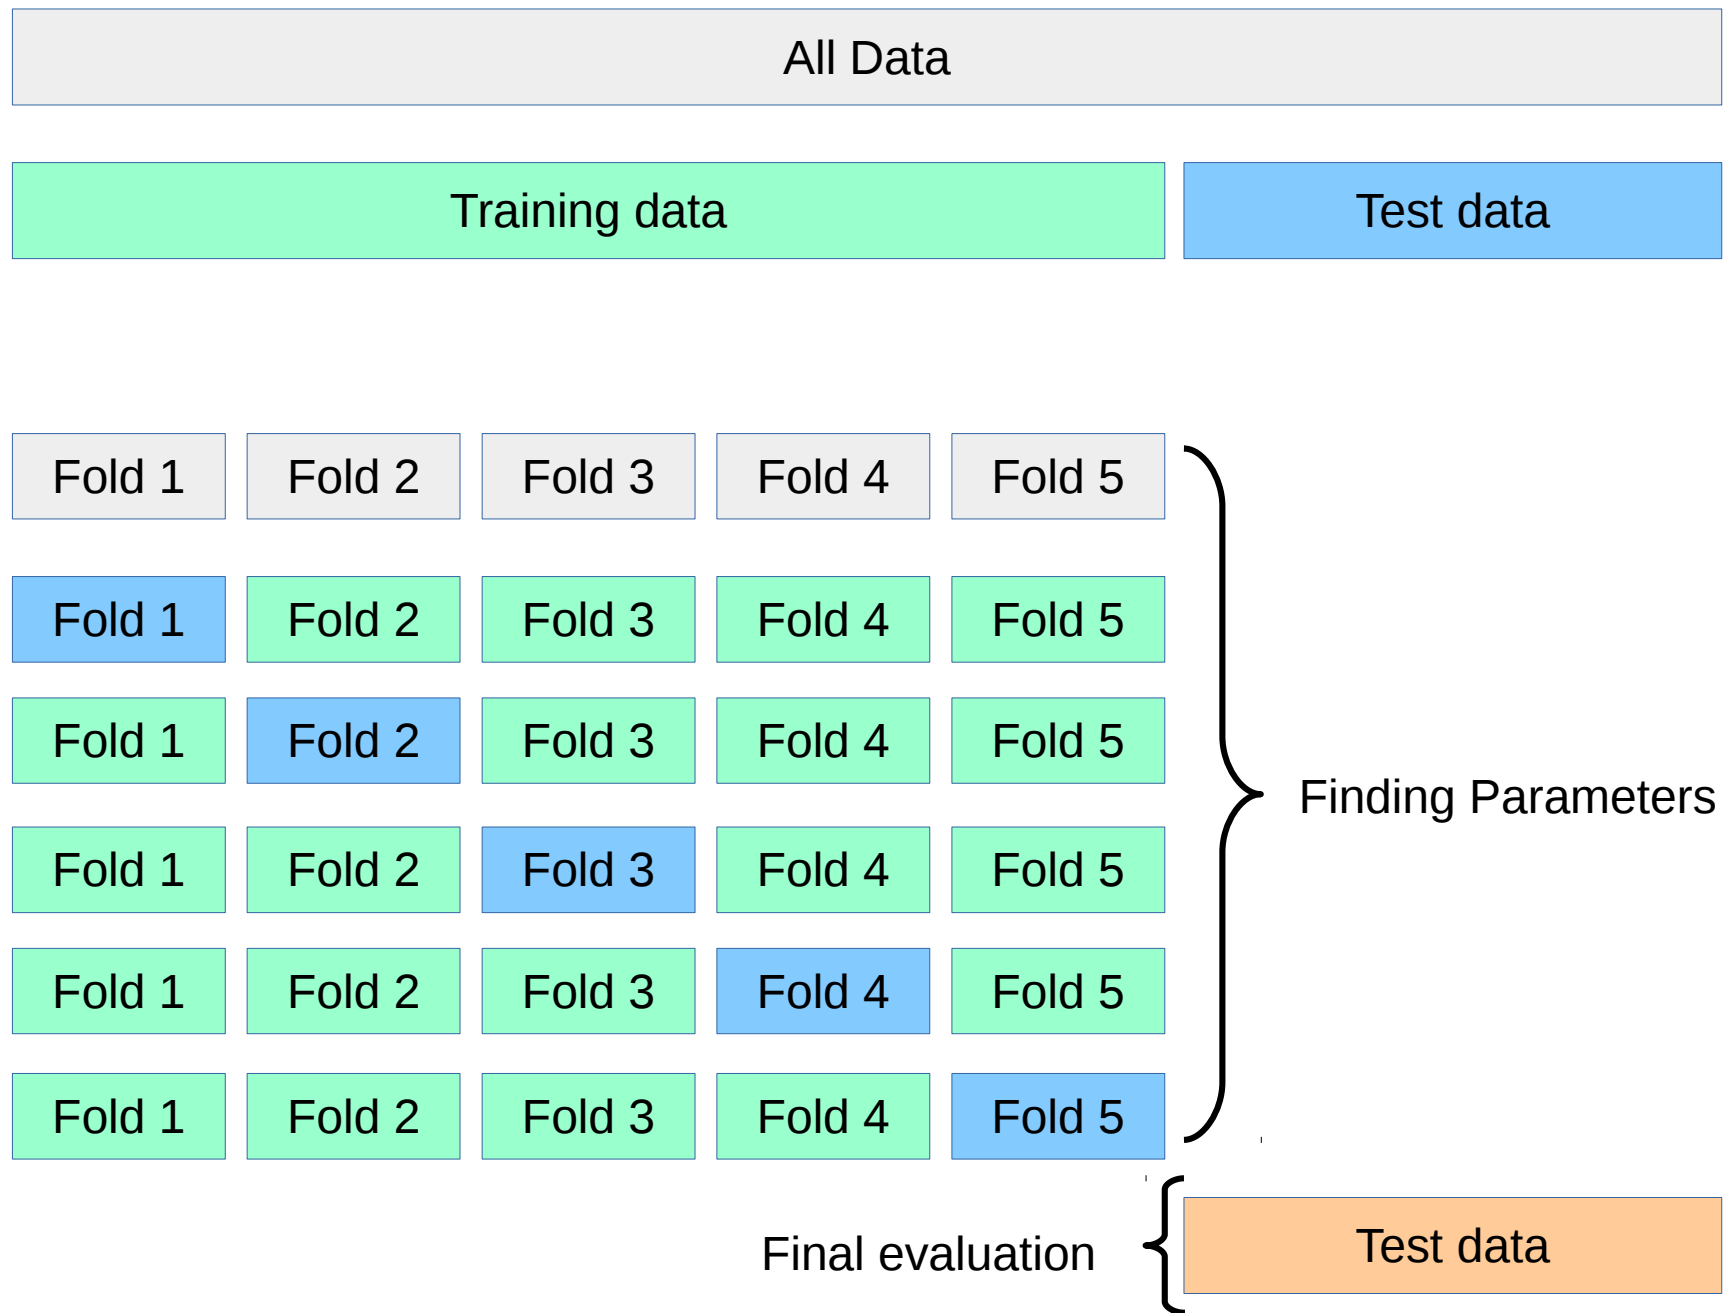
Split 2   Fold 1   Fold 2   Fold 3   Fold 4   Fold 5

Split 3   Fold 1   Fold 2   Fold 3   Fold 4   Fold 5

Split 4   Fold 1   Fold 2   Fold 3   Fold 4   Fold 5

Split 5   Fold 1   Fold 2   Fold 3   Fold 4   Fold 5

Test data



# Cross -Validated Grid Search

```
from sklearn.grid_search import GridSearchCV
from sklearn.cross_validation import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y)

param_grid = {'C': 10. ** np.arange(-3, 3),
              'gamma': 10. ** np.arange(-3, 3)}
grid = GridSearchCV(SVC(), param_grid=param_grid)
grid.fit(X_train, y_train)
grid.predict(X_test)
grid.score(X_test, y_test)
```

## Sample application: Sentiment Analysis

# IMDB Movie Reviews Data

## **Review:**

One of the worst movies I've ever rented. Sorry it had one of my favorite actors on it (Travolta) in a nonsense role. In fact, anything made sense in this movie.

Who can say there was true love between Eddy and Maureen?  
Don't you remember the beginning of the movie ?

Is she so lovely? Ask her daughters. I don't think so.

**Label:** negative

**Training data:** 12500 positive, 12500 negative

# Bag Of Word Representations

`CountVectorizer / TfidfVectorizer`

# Bag Of Word Representations

`CountVectorizer / TfidfVectorizer`

`"This is how you get ants."`



# Bag Of Word Representations

CountVectorizer / TfidfVectorizer

"This is how you get ants."

tokenizer



['this', 'is', 'how', 'you', 'get', 'ants']

# Bag Of Word Representations

CountVectorizer / TfidfVectorizer

"This is how you get ants."

tokenizer

↓  
['this', 'is', 'how', 'you', 'get', 'ants']

↓  
Build a vocabulary over all documents

↓  
['aardvak', 'amsterdam', 'ants', ... 'you', 'your', 'zyxst']

# Bag Of Word Representations

CountVectorizer / TfidfVectorizer

"This is how you get ants."

tokenizer

['this', 'is', 'how', 'you', 'get', 'ants']

Build a vocabulary over all documents

['aardvak', 'amsterdam', 'ants', ... 'you', 'your', 'zyxst']

Sparse matrix encoding

aardvak	ants	get	you	zyxst
[0, ..., 0, 1, 0, ... , 0, 1 , 0, ..., 0, 1, 0, .....	0]			

# N-grams (unigrams and bigrams)

`CountVectorizer / TfidfVectorizer`

# N-grams (unigrams and bigrams)

`CountVectorizer / TfidfVectorizer`

`"This is how you get ants."`

# N-grams (unigrams and bigrams)

`CountVectorizer` / `TfidfVectorizer`

"This is how you get ants."

Unigram tokenizer



`['this', 'is', 'how', 'you', 'get', 'ants']`

# N-grams (unigrams and bigrams)

CountVectorizer / TfidfVectorizer

"This is how you get ants."

Unigram tokenizer

↓  
['this', 'is', 'how', 'you', 'get', 'ants']

"This is how you get ants."


Bigram tokenizer


↓  
['this is', 'is how', 'how you', 'you get', 'get ants']


# Notebook Working With Text Data





# Pipelines









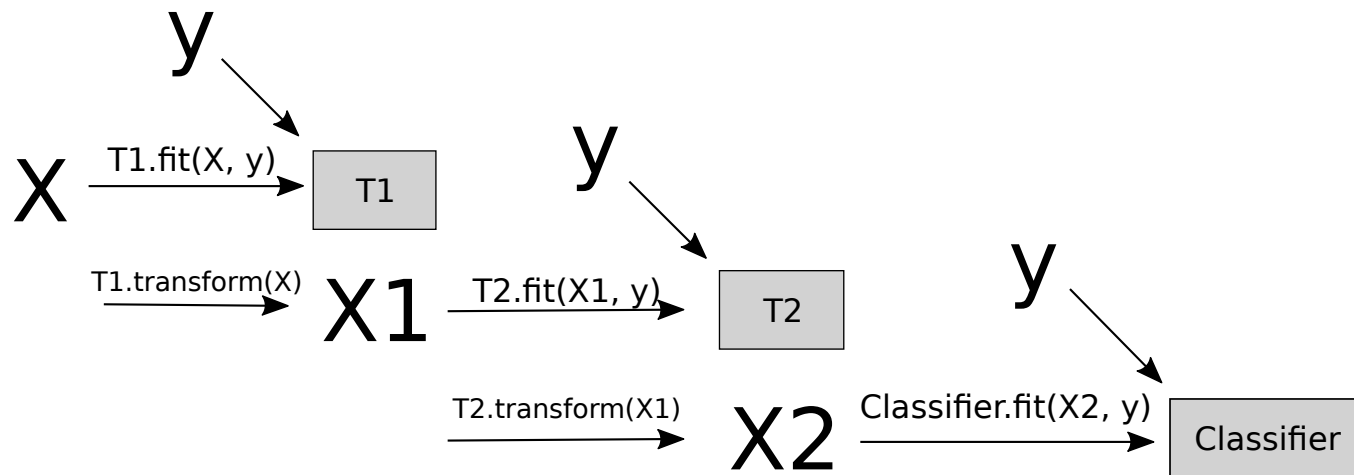


# Pipelines


```
pipe = make_pipeline(T1(), T2(), Classifier())
```



```
pipe.fit(X, y)
```



```
pipe.predict(X')
```



# Pipelines


```
from sklearn.pipeline import make_pipeline  
  
pipe = make_pipeline(StandardScaler(), SVC())  
pipe.fit(X_train, y_train)  
pipe.predict(X_test)
```




# Randomized Parameter Search

# Randomized Parameter Search

Grid Layout




Random Layout




# Randomized Parameter Search

Grid Layout



Random Layout



Step-size free for continuous parameters  
Decouples runtime from search-space size  
Robust against irrelevant parameters

# Randomized Parameter Search


```
params = {'featureunion__countvectorizer-1__ngram_range':  
          [(1, 3), (1, 5), (2, 5)],  
          'featureunion__countvectorizer-2__ngram_range':  
          [(1, 1), (1, 2), (2, 2)],  
          'linearsvc__C': 10. ** np.arange(-3, 3)}
```

# Randomized Parameter Search

```
params = {'featureunion__countvectorizer-1__ngram_range':  
          [(1, 3), (1, 5), (2, 5)],  
          'featureunion__countvectorizer-2__ngram_range':  
          [(1, 1), (1, 2), (2, 2)],  
          'linearsvc__C': expon()}
```

# Randomized Parameter Search

```
params = {'featureunion__countvectorizer-1__ngram_range':  
          [(1, 3), (1, 5), (2, 5)],  
          'featureunion__countvectorizer-2__ngram_range':  
          [(1, 1), (1, 2), (2, 2)],  
          'linearsvc__C': expon()}
```




```
rs = RandomizedSearchCV(text_pipe,  
                        param_distributions=param_distributins, n_iter=50)
```

# Randomized Parameter Search

- Always use distributions for continuous variables.
- Don't use for low dimensional spaces.

# GP based parameter optimization (coming soon)



From Eric Brochu, Vlad M. Cora and Nando de Freitas



# Scoring Functions

GridSeachCV  
RandomizedSearchCV  
cross\_val\_score  
...CV

Default:  
Accuracy (classification)  
R2 (regression)

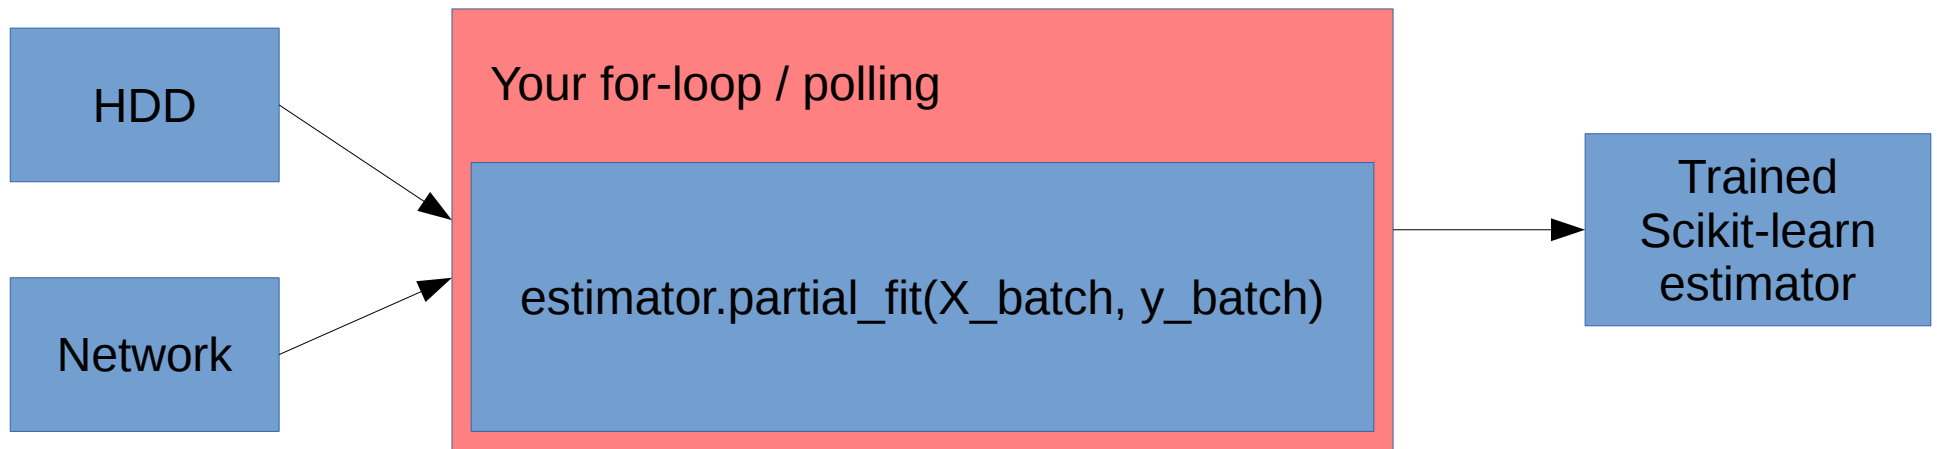
# Notebook scoring metrics

# Out of Core Learning

# Three regimes of data

- Fits in RAM
- Fits on a Hard Drive
- Doesn't fit on a single PC

	vCPU	ECU	Memory (GiB)	Instance Storage (GB)	Linux/UNIX Usage
<b>Memory Optimized - Current Generation</b>					
r3.large	2	6.5	15	1 x 32 SSD	\$0.195 per Hour
r3.xlarge	4	13	30.5	1 x 80 SSD	\$0.39 per Hour
r3.2xlarge	8	26	61	1 x 160 SSD	\$0.78 per Hour
r3.4xlarge	16	52	122	1 x 320 SSD	\$1.56 per Hour
r3.8xlarge	32	104	244	2 x 320 SSD	\$3.12 per Hour
<b>Storage Optimized - Current Generation</b>					
i2.xlarge	4	14	30.5	1 x 800 SSD	\$0.938 per Hour
i2.2xlarge	8	27	61	2 x 800 SSD	\$1.876 per Hour
i2.4xlarge	16	53	122	4 x 800 SSD	\$3.751 per Hour
i2.8xlarge	32	104	244	8 x 800 SSD	\$7.502 per Hour



# Supported Algorithms

- All `SGDClassifier` derivatives
- Naive Bayes
- `MinibatchKMeans`
- `IncrementalPCA`
- `MiniBatchDictionaryLearning`
- `MultilayerPerceptron` (dev branch)
- `Scalers`



# Out of Core Learning

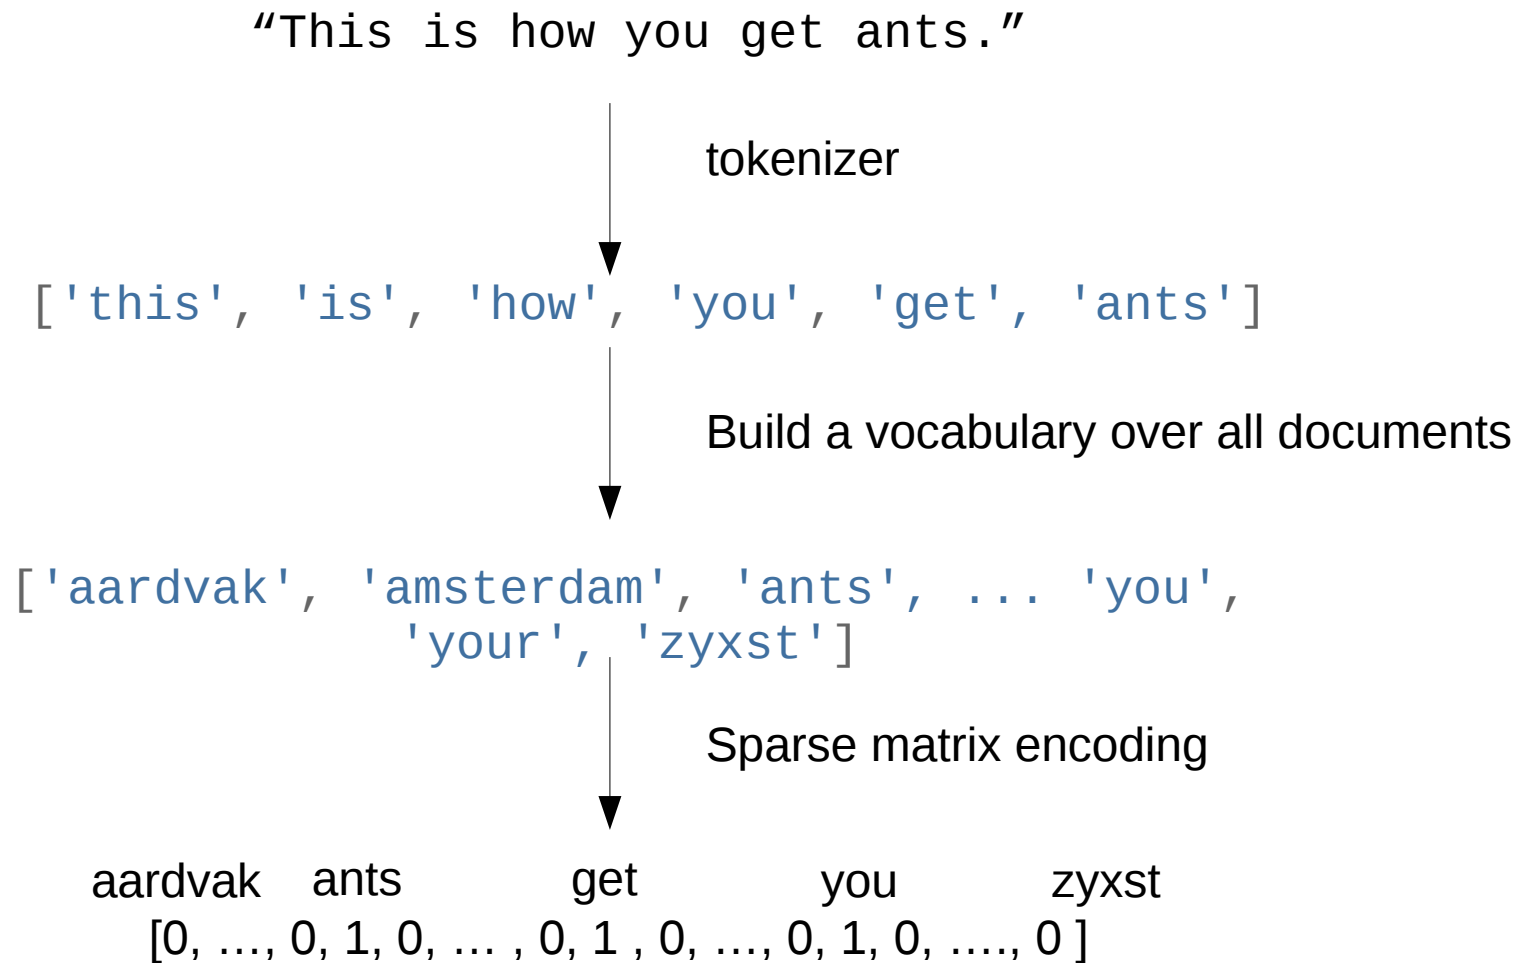
```
sgd = SGDClassifier()

for i in range(9):
    X_batch, y_batch = cPickle.load(open("batch_%02d" % i))
    sgd.partial_fit(X_batch, y_batch, classes=range(10))
```

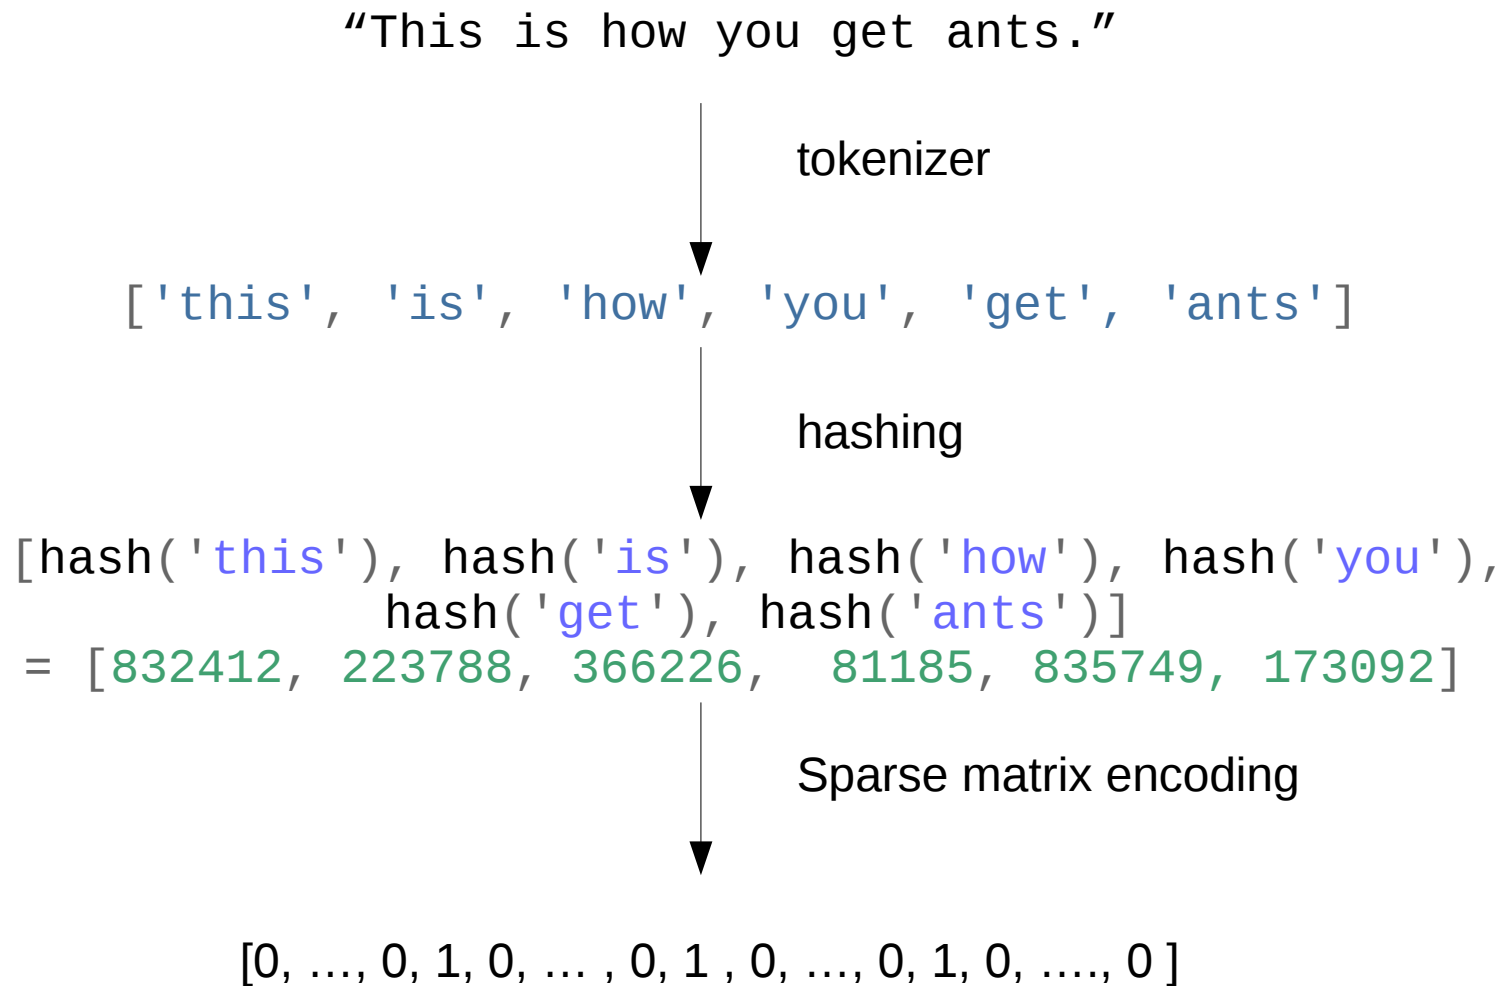
Possibly go over the data multiple times.

# The hashing trick for text data

# Text Classification: Bag Of Word



# Text Classification: Hashing Trick



# Text Classification: Hashing Trick

```
sgd = SGDClassifier()
hashing_vectorizer = HashingVectorizer()

for batch_name in glob("*.pickle"):
    with open(batch_name) as f:
        text_batch, y_batch = pickle.load(batch_name)

    X_batch = hashing_vectorizer.transform(text_batch)
    sgd.partial_fit(X_batch, y_batch, classes=[0, 1])
```

How (and why) to build your own estimator

# Why?

GridSearchCV  
cross\_val\_score  
Pipeline

# How

- “fit” method
- set\_params and get\_params (or inherit)
- Run check\_estimator

See the “build your own estimator” docs!




# Notebook Building your own estimator

What's new in 0.17

# Latent Dirichlet Allocation

using online variational inference



Topic #0:

government people mr law gun state president states public use right  
rights national new control american security encryption health  
united

Topic #1:

drive card disk bit scsi use mac memory thanks pc does video hard  
speed apple problem used data monitor software

Topic #2:

said people armenian armenians turkish did saw went came women  
killed children turkey told dead didn left started greek war

Topic #3:

year good just time game car team years like think don got new play  
games ago did season better ll

Topic #4:

10 00 15 25 12 11 20 14 17 16 db 13 18 24 30 19 27 50 21 40

Topic #5:

windows window program version file dos use files available display  
server using application set edu motif package code ms software

Topic #6:

edu file space com information mail data send available program ftp  
email entry info list output nasa address anonymous internet

Topic #7:

ax max b8f g9v a86 pl 145 1d9 0t 34u 1t 3t giz bhj wm 2di 75u 2tm  
bxn 7ey

Topic #8:


god people jesus believe does say think israel christian true life jews  
did bible don just know world way church

Topic #9:

don know like just think ve want does use good people key time way  
make problem really work say need

By Chyi-Kwei Yau, based on code by Matt Hoffman

# SAG for Logistic Regression and Ridge Regression



By Danny Sullivan and Tom Dupre la Tour

# Coordinate Descent Solver for Non-Negative Matrix Factorization

Topics in NMF model:

Topic #0:

don people just like think know time good right ve make say want did really way new use going said

Topic #1:

windows file dos files window program use running using version ms problem server pc screen ftp run application os software

Topic #2:

god jesus bible christ faith believe christians christian heaven sin hell life church truth lord say belief does existence man

Topic #3:

geb dsl n3jxp chastity cadre shameful pitt intellect skepticism surrender gordon banks soon edu lyme blood weight patients medical probably

Topic #4:

key chip encryption clipper keys escrow government algorithm secure security encrypted public des nsa enforcement bit privacy law secret use

Topic #5:

drive scsi ide drives disk hard controller floppy hd cd mac boot rom cable internal tape bus seagate bios quantum

Topic #6:

game team games players year hockey season play win league teams nhl baseball player detroit toronto runs pitching best playoffs

Topic #7:

thanks mail does know advance hi info looking anybody address appreciated help email information send ftp post interested list appreciate


Topic #8:

card video monitor vga bus drivers cards color driver ram ati mode memory isa graphics vesa pc vlb diamond bit

Topic #9:

00 sale 50 shipping 20 10 price 15 new 25 30 dos offer condition 40 cover asking 75 interested 01

# Barnes-Hut Approximation for T-SNE manifold learning



# FunctionTransformer

```
>>> import numpy as np
>>> from sklearn.preprocessing import FunctionTransformer
>>> transformer = FunctionTransformer(np.log1p)
>>> X = np.array([[0, 1], [2, 3]])
>>> transformer.transform(X)
array([[ 0.          ,  0.69314718],
       [ 1.09861229,  1.38629436]])
```


# VotingClassifier

```
clf1 = LogisticRegression()  
clf2 = RandomForestClassifier()  
clf3 = GaussianNB()  
  
ecf = VotingClassifier(  
    estimators=[('lr', clf1), ('rf', clf2), ('gbn', clf3)],  
    voting="hard")
```



What the future will bring (0.18)


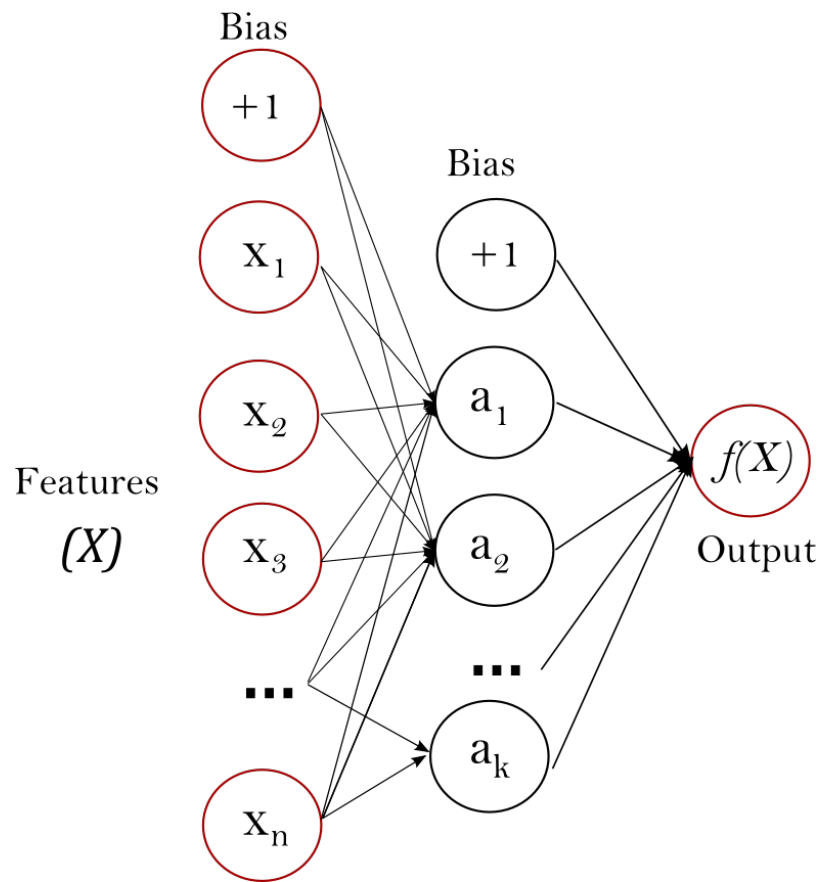
# Gaussian Process Rewrite



```
34.4**2 * RBF(length_scale=41.8)
+ 3.27**2 * RBF(length_scale=180)
    * ExpSineSquared(length_scale=1.44, periodicity=1)
+ 0.446**2 * RationalQuadratic(alpha=17.7, length_scale=0.957)
+ 0.197**2 * RBF(length_scale=0.138) +
WhiteKernel(noise_level=0.0336)
```

By Jan Hendrik Metzen.

# Neural Networks



By Jiyuan Qian and Issam Laradji

# Improved Cross-Validation

current

```
>>> import numpy as np
>>> from sklearn.cross_validation import KFold

>>> kf = KFold(4, n_folds=2)
>>> for train, test in kf:
...     print("%s %s" % (train, test))
[2 3] [0 1]
[0 1] [2 3]
```


future

```
>>> import numpy as np
>>> from sklearn.model_selection import KFold

>>> X = ["a", "b", "c", "d"]
>>> kf = KFold(n_folds=2)
>>> for train, test in kf.split(X):
...     print("%s %s" % (train, test))
[2 3] [0 1]
[0 1] [2 3]
```

```
gs = GridSearchCV(LinearSVC(random_state=0), param_grid={'C': [1, 10]},
                  cv=inner_cv)
cross_val_score(gs, X=X, y=y, labels=labels, cv=outer_cv,
                fit_params={'labels': labels})
```

# Faster PCA



By Giorgio Patrini

O'REILLY®




# Introduction to Machine Learning with Python

A GUIDE FOR DATA SCIENTISTS

Andreas C. Müller & Sarah Guido

Release June 2016

Hellbender



Release June 2016


# Thank you!



@amuellerm1



@amueller



importamueller@gmail.com



<http://amueller.github.io>